

— HONEST STATUS REPORT • SUPERFECTA HELP • 2026-04-15

Superfecta Project Report

A clean, report-safe read on what actually improved, what deserves paper-trade trust now, and what still does *not* deserve hype.

Phase 7 = primary paper baseline

Phase 8 = shadow challenger

Selective rules = only live-worthy family

Main conclusion: the project got meaningfully better through **validation cleanup**, **decision clarity**, and **operational hardening**, not through fake “overnight edge” claims.

BEST CLEAN GAIN

+7.96 pts

Walk-forward selector ROI improved from +22.46% to +30.42% with sqrt-dampened scoring.

PRIMARY BASELINE

Phase 7 live portfolio

+38.68%

SHADOW CHALLENGER

Phase 8 frozen portfolio

+21.45%

LIVE PEAK TODAY

Best demo EV seen

+23.24%

ALERT BAR

Watcher threshold

+30.00%

HOLDOUT WINNER

+38.68%

Phase 7 holdout ROI on **175 races**, currently the strongest practical paper baseline.

SELECTOR UPGRADE

+30.42%

Best validated walk-forward selector result after scoring cleanup, up from +22.46%.

OPERATIONAL SHIFT

Paperable

Daily paper-trade, settlement, forward-check, lane monitor, and preflight note path are now in place.

LIVE STATUS

No +30 yet

Demo/watcher lane is working, but today's best live EV stayed below the real alert threshold.

WHAT ACTUALLY GOT BETTER

The biggest real improvements

This repo is stronger now because the evidence is cleaner, the operating path is real, and the live/demo lane no longer gets confused with production.

1

Validation got sharper

The strongest clean numeric gain is the selector improvement from **+22.46%** to **+30.42%**, using frozen artifacts and no new rule mining.

2

Portfolio hierarchy got clearer

Phase 7 remains the live paper baseline. Phase 8 stays useful, but only as a shadow challenger until it beats Phase 7 on better forward evidence.

3

Operations got real

The project now has daily paper-trade runners, settlement tooling, forward checks, lane monitors, and preflight context for quiet race days.

BEFORE

Strong but messy

Good research, too much drift, too many files saying half of the story.

NOW

Defensible + usable

Cleaner deployment logic, clearer ranking of what is paper / shadow / benchmark / skip.

NET EFFECT


Less bullshit

Better odds of saying the true thing quickly, especially under pressure or while showing work to Cole.

QUANTIFIED CHANGE

Numbers that matter

These are the metrics worth saying out loud. They hold up better than the older shiny-but-looser headline stats.



The cleanest sentence in the whole project is still this: **Phase 7 beats Phase 8 on the current holdout, even though Phase 8 looks prettier on paper.**

PORTFOLIO CALL

What to run, what to watch, what to cite

The live choice, the shadow choice, and the benchmark choice are not the same thing anymore. That’s good.

APPROACH	ROLE	HOLDOUT	WALK-FORWARD	READ
Phase 7 live portfolio	PAPER NOW	+38.68% • 175 races	+31.34% • 806 races	Strongest current paper baseline. BEL is dormant, so this is effectively the OP+CD lane.
Phase 8 frozen portfolio	SHADOW ONLY	+21.45% • 118 races	+55.04% • 625 races	Useful challenger, but weaker on the current holdout despite looking hotter in broader mined views.
Train-only yearly selector	BENCHMARK	+14.36% • 65 races	+22.46% • 470 races	Still the honesty yardstick, just not the best operating

APPROACH ROLE HOLDOUT WALK-FORWARD READ

recipe for
daily use.

PRIMARY CALL

Phase 7 first

Use the Phase 7 live basket as the real paper-trade default because it owns the best current holdout evidence.

CHALLENGE LANE

Phase 8 in shadow

Log it, compare it, but do not promote it just because it has more moving pieces or a prettier walk-forward headline.

METHOD-FAMILY TRIAGE

Which family deserves trust?

This may be the most useful cleanup in the whole repo: what belongs in the real decision path versus what should stay benchmark or research.

METHOD FAMILY	STATUS	BEST HONEST EVIDENCE	READ
Selective rule path	PAPER NOW	+38.68% holdout • +31.34% WF	Only family with positive current frozen holdout evidence and a practical live paper-trade path.
Harville-ranked probabilities	BENCHMARK ONLY	-24.05% ROI • 90,004 races	Useful structural baseline, but the takeout still wins if you stop at Harville ranking alone.
XGBoost residual correction	RESEARCH ONLY	-24.16% ML betting ROI • 16,724 races	Improves payout prediction a bit, but the downstream EV betting behavior barely

METHOD FAMILY	STATUS	BEST HONEST EVIDENCE	READ
			changes. Not enough to promote.

RULE STACK

Current rule hierarchy

These are the rules worth naming first when somebody asks what the project actually trusts.

RULE	TIER	HOLDOUT	WF
OP_DURABLE_K7	ANCHOR	+22.9% • 115	7/10 selected
CD_CORE_K8	PAPER	+56.0% • 60	1/10 selected
OP_REFINED_K7	WATCH	+51.4% • 49	2/10 selected

OPERATIONAL HARDENING

Why the project feels more real now

The strongest practical gain is not a headline stat, it's that the daily operating path stopped being hand-wavy.

DAILY

Primary + shadow runners

Separate paper lanes make it obvious what is the real baseline versus what is only being watched.

AFTER RACES

Settlement sync + helper

There is now a clean place to record actual outcomes and returns instead of hand-editing ledgers and hoping nothing drifts.

VALIDATION

Forward checks + lane monitors

Forward evidence can now be read as *no data*, *too early*, *running cold*, or *within noise* instead of being guessed from raw CSVs.

QUIET DAYS

Preflight note

Quiet days now explain whether OP/CD simply were not racing, instead of making the project look broken or low-signal by accident.

LIVE / DEMO LANE

Today’s live reality

The live model lane works. The watcher works. The threshold discipline is the point.

WORKING NOW

Separate demo lane

`demo\_live\_predictions.py` now handles real available cards without pretending the production OP/CD basket changed.

- Best-live selection mode
- PLAY / FLOW / PASS labeling
- Candidate leaderboard output
- Cache hardening + latest JSON snapshots
- Silent watcher unless a true +30% EV spot appears

TODAY’S BEST

Laurel R9 • +23.24%

Encouraging, but still below the alert bar. That is why the watcher stays quiet.



WHAT STILL DOES NOT CLEAR

Do not over-claim these

This project got better partly because some tempting stories were explicitly rejected.

**Do not say Phase 8 replaced Phase 7.**

It didn't. Phase 8 is still shadow-only until it wins on better forward evidence.

**Do not say ML fixed the betting case.**

Horse-history features improved prediction quality, but downstream EV decisions barely changed.

**Do not claim a +30 live opportunity happened today.**

Best seen was +23.24%, which is below the real alert bar.

**Do say the project is much more defensible now.**

The report hierarchy, paper-trade path, and live/demo separation are all materially better than before.

FINAL STANCE

The clean recommendation

If somebody only remembers one paragraph from this report, it should be this one.

The project improved meaningfully, but in the right way: **better validation, better decision hygiene, better paper-trade operations**, and a **working live/demo lane** that does not fake production readiness. The strongest current deployment stance is still conservative: **paper trade the selective rule path, keep Phase 7 as the primary baseline, keep Phase 8 as shadow-only, and keep Harville/XGBoost out of the live betting decision path.**

Built from the current Superfecta evidence stack on **2026-04-15**.

Core inputs: compare_main_approaches, forward_evidence_scorecard, portfolio decision card, method-family decision card, selector experiment, sample-size experiment, overnight progress, and the live demo watcher outputs.